

Python and LaTeX: Assignment 2

Matthew DuBois

October 18, 2024

1 Introduction

The goal of this assignment was to implement a linear classifier using both SVM and softmax functions. After these functions were implemented, we also used gradient regularization and gradient descent to optimize the function.

2 Python Code

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 N = 100 # number of points per class
5 D = 2 # dimensionality
6 K = 3 # number of classes
7 X = np.zeros((N*K,D)) # data matrix (each row = single
   # example)
8 y = np.zeros(N*K, dtype='uint8') # class labels
9 for j in range(K):
10     ix = range(N*j,N*(j+1))
11     r = np.linspace(0.0,1,N) # radius
12     t = np.linspace(j*4,(j+1)*4,N) + np.random.randn(N)
   # *0.2 # theta
13     X[ix] = np.c_[r*np.sin(t), r*np.cos(t)]
14     y[ix] = j
15 # lets visualize the data:
16 plt.scatter(X[:, 0], X[:, 1], c=y, s=40, cmap=plt.cm.
   Spectral)
17 plt.show()
18 W = np.random.randn(2, 3)
19 b = np.random.randn(3)
20 def score_function(X, W, b):
21     scores = np.dot(X, W) + b
22     return scores
```

```

23 scores = score_function(X, W, b)
24 print("First 1 row of scores:\n", scores[:1])

1 import numpy as np
2
3
4 def svm_loss(scores, y):
5     """
6     Calculate the multiclass SVM loss for a single
7     example.
8
9     :param scores: numpy array of scores for each
10                    class (output of the classifier)
11     :param y: integer, the correct class index
12     :return: float, the loss for the input sample
13     """
14     correct_class_score = scores[y]
15     margin = 1
16     losses = np.maximum(0, scores -
17                          correct_class_score + margin)
18     losses[y] = 0
19     return np.sum(losses)
20
21 scores = np.array([0.2, 0.4, 0.1])
22 y = 1
23 print("SVM Loss:", svm_loss(scores, y))
24
25 def softmax_loss(scores, y):
26     """
27     Calculate the cross-entropy loss using softmax for
28     a single example.
29
30     :param scores: numpy array of scores from the
31                    model for each class
32     :param y: integer, the correct class index
33     :return: float, the loss for the input sample
34     """
35     shift_scores = scores - np.max(scores)
36     exp_scores = np.exp(shift_scores)
37     softmax_probabilities = exp_scores / np.sum(
38         exp_scores)
39     return -np.log(softmax_probabilities[y])

```

```

37 scores = np.array([0.2, 0.4, 0.1])
38 y = 1
39 print("Softmax Loss:", softmax_loss(scores, y))
40
1 def softmax_loss_with_regularization(scores, y, W,
2   lambda_reg):
3     """
4     Calculate the cross-entropy loss using softmax for
5     a single example with L2 regularization.
6
7     :param scores: numpy array of scores from the
8     model for each class
9     :param y: integer, the correct class index
10    :param W: numpy array, weight matrix of the
11    classifier
12    :param lambda_reg: float, regularization strength
13    :return: float, the regularized loss for the input
14    sample
15    """
16
17    shift_scores = scores - np.max(scores)
18    exp_scores = np.exp(shift_scores)
19    softmax_probabilities = exp_scores / np.sum(
20        exp_scores)
21    cross_entropy_loss = -np.log(softmax_probabilities
22        [y])
23
24    l2_reg = lambda_reg * np.sum(W**2)
25
26    return cross_entropy_loss + l2_reg
27
28 scores = np.array([0.2, 0.4, 0.1])
29 y = 1
30 W = np.array([[0.1, 0.2, 0.3], [0.4, 0.5, 0.6]])
31 lambda_reg = 0.1
32
33 print("Softmax Loss with Regularization:",
34     softmax_loss_with_regularization(scores, y, W,
35     lambda_reg))
36
1 import numpy as np
2
3 def compute_softmax_loss_and_gradient(W, X, y,

```

```

lambda_reg):
4     num_examples = X.shape[0]  # Number of training
      examples
5     scores = np.dot(X, W)
6     shift_scores = scores - np.max(scores, axis=1,
      keepdims=True)  # Numeric stability
7     exp_scores = np.exp(shift_scores)
8     softmax_probabilities = exp_scores / np.sum(
      exp_scores, axis=1, keepdims=True)
9
10    correct_log_probs = -np.log(softmax_probabilities[
      range(num_examples), y])
11    data_loss = np.sum(correct_log_probs) /
      num_examples
12    reg_loss = 0.5 * lambda_reg * np.sum(W * W)
13    loss = data_loss + reg_loss
14
15    dscores = softmax_probabilities
16    dscores[range(num_examples), y] -= 1
17    dscores /= num_examples
18    dW = np.dot(X.T, dscores)
19    dW += lambda_reg * W
20
21    return loss, dW
22
23
24    num_examples = 300
25    dim = 5
26    num_classes = 3
27
28    np.random.seed(42)
29    X = np.random.randn(num_examples, dim)
30    y = np.random.randint(num_classes, size=num_examples)
31    learning_rate = 1e-2
32    num_iterations = 200
33    lambda_reg = 0.1
34
35
36    W = 0.001 * np.random.randn(dim, num_classes)
37
38
39    for i in range(num_iterations):
40        loss, grad = compute_softmax_loss_and_gradient(W,
            X, y, lambda_reg)
41        W -= learning_rate * grad  # Update weights
42

```

```

43         if i % 10 == 0:
44             print(f'Iteration {i}, loss: {loss}')
45
46     print("Final loss:", loss)
47     print("Optimized weights:\n", W)

```

3 Methodology

3.1 Question 2

For question two, we implemented a linear score function. Given that this is a spiral data set, a linear score function would not be ideal for classification. This resulted in the low confidence scores that were outputs of the function that are posted in the section below.

3.2 Question 3

Both the hinge loss and softmax loss functions were used to train the model. The hinge loss is used to create better separation between the different classes. This works by taking the score of the correct class, and checking the other scores against a margin. If the margin is too close, it will be added to total loss. The addition of the softmax function was to increase the confidence in the model. This works by turning the scores into probabilities. A high probability results in a small loss, whereas a low probability results in a high loss.

3.3 Question 4

The addition of regularization helps to smooth out the model and make it simpler. This is done by adding a penalty to the loss function based on the weights. With a penalty assigned to the weights, this helps to prevent the model from assigning too much weight to one component.

3.4 Question 5

The gradient descent was used to optimize the loss function. This requires first calculating the gradient descent for the loss function, then adjusting the weights in the opposite direction to decrease the loss. The learning rate determines this adjustment and after a number of iterations, the models performance will improve.

4 Results

4.1 q1

First 1 row of scores: $[-0.06769397 \ 0.54871401 \ 1.73663602]$

4.2 q2

SVM Loss: 1.5 Softmax Loss: 0.93983106084446

4.3 q3

Softmax Loss with Regularization: 1.03083106084446

4.4 q4

Iteration 0, loss: 1.0988139923881797 Iteration 10, loss: 1.0980923175227943
Iteration 20, loss: 1.0974293545604854 Iteration 30, loss: 1.0968202494405201
Iteration 40, loss: 1.096260553941703 Iteration 50, loss: 1.0957461917901279
Iteration 60, loss: 1.0952734275141656 Iteration 70, loss: 1.0948388378457865
Iteration 80, loss: 1.0944392854773455 Iteration 90, loss: 1.094071894993488
Iteration 100, loss: 1.0937340308086125 Iteration 110, loss: 1.0934232769511143
Iteration 120, loss: 1.093137418546245 Iteration 130, loss: 1.0928744248597562
Iteration 140, loss: 1.092632433774427 Iteration 150, loss: 1.0924097375810833
Iteration 160, loss: 1.0922047699747033 Iteration 170, loss: 1.0920160941547077
Iteration 180, loss: 1.0918423919365023 Iteration 190, loss: 1.0916824537887841
Final loss: 1.0915493584063678 Optimized weights: [[0.04230676 -0.03577001 -
0.00534934] [-0.00521211 0.02514545 -0.02226465] [-0.04790917 -0.00740088 0.05553504]
[0.01025628 -0.00743917 -0.00209458] [-0.03815491 0.04351046 -0.00617217]]

5 Conclusion

This assignment tested our knowledge of linear classifiers and gradient descents. After generating a spiral dataset, an SVM and softmax linear classifier were implemented in the function to help train the model. After training, regularization was implemented which helps to prevent the over weighting of a single component. Finally, a gradient descent was used to optimize the function and find the resulting weights.